Lingzhi Zhao

Ph.D. Candidate in Computer Science | University of Illinois Urbana-Champaign
+1 (217) 979-9381 | lz26@illinois.edu | Website | LinkedIn

RESEARCH INTERESTS

My research bridges **XR systems** and **multimodal AI**. I am interested in advancing **XR systems** by optimizing **system performance** and **integrating VLMs** for better interaction experience.

EDUCATION

University of Illinois Urbana-Champaign · Ph.D. in Computer Science — Aug. 2022 – May 2027
Advisor: Prof. Klara Nahrstedt

Shanghai Jiao Tong University · M.S. in Electrical and Computer Engineering — 2019 – 2022

Shanghai University · B.S. in Electrical and Computer Engineering — 2015 – 2019

TECHNICAL SKILLS

ML	PyTorch, Transformers, VLM/MLLM, SFT, on-device inference
Programming	Python, C++, C#, Swift, Java, Rust
Tools	FFmpeg, NVCodec, LLMFarm, Wireshark, iperf3, tc, Git

RESEARCH & WORK EXPERIENCE

University of Illinois Urbana-Champaign, MONET Lab — Aug. 2022 – Present

Graduate Research Assistant · Advisor: Prof. Klara Nahrstedt

AquaVLM: VLM-Powered VR for Underwater Communication *[ACM UbiComp 2027] [Website]*

- Built an **end-to-end interactive VR application** for underwater communication, using **on-device VLM inference** over **multimodal inputs** to enable divers to generate and exchange context-aware messages.
- Fine-tuned **MobileVLM-V2 (3B) via SFT** on an **85K-sample multimodal dataset** constructed with GPT-4o few-shot prompting to support context-aware message generation.
- Designed a **hierarchical message-generation** pipeline, achieving **>80% user-intent alignment** while reducing inference latency by **60%–70%**.

GenStream: Generative Codecs for Real-Time Video Streaming

- Optimized a generative codec for **real-time video streaming**, achieving **50%–78% bandwidth savings** over H.264/H.265 and prior neural codecs at comparable quality.
- Designed **content-aware token adaptation** for changing network conditions, improving perceptual quality (LPIPS) by **38%–59%** at matched bitrates and loss patterns versus state-of-the-art streaming systems.

AquaScope: Generative Compression for Mobile Image Transmission *[ACM MobiSys 2026] [Website]*

- Optimized visual tokenizers (TiTok, VQGAN) through **token clustering and fine-tuning** for **low-bandwidth, error-resilient image transmission**.
- Integrated the learned codec into a mobile acoustic prototype, improving **bandwidth efficiency** by **40%** and reducing average **BER** from **19% to below 2%**; outperformed PNG, JPEG, and classic neural compression.

ByteDance Inc., San Diego, CA — May 2024 – Aug. 2025

Research Intern, Multimedia Lab · Manager: Dr. Shu Shi

MARS: Motion-Aware ABR for Real-Time VR *[Submitted to ACM MMSys 2027]*

- Developed **perceptually guided bitrate adaptation for real-time VR**, reducing bandwidth usage by **15%** without degrading quality of experience (QoE), validated in a user study with diverse VR applications.
- Designed a **kinematics-based state machine** to synchronize bitrate adaptation with **head-motion phases**; integrated motion-vector information and support for CBR, VBR, and CQP rate control for the system.

Trinity: Low-Latency Cloud VR Gaming *[ACM MMSys 2026] [Website]*

- Prototyped an **end-to-end cloud VR gaming system**, reducing bandwidth cost by **60%** and GPU utilization by **40%** while maintaining QoE comparable to local streaming.
- Characterized **latency tolerance across head, hand, and body motion** using an open-source first-person shooter VR game.
- Designed **delayed rendering for hand motion** and **partial-panorama rendering for head motion** to account for action-dependent latency sensitivity.

TikTok, San Jose, CA May 2026 – Present

Machine Learning Engineer Intern · Manager: Prof. Feng Qian

MLLM-Based Perceptual Quality Assessment for Short Videos *[Submitted to ICLR 2027]*

- Built a **short-video transcoding quality benchmark** with **377K subjective ratings**, spanning source resolutions, transcoding settings, and mobile devices; analyzed MLLM scorers' failure modes.
- Post-trained a **context-aware MLLM with motion–artifact fusion** via SFT, achieving state-of-the-art quality prediction on the new benchmark and three public datasets.

Dynamic Ad Object Insertion for UGC Videos *[Manuscript in preparation]*

- Built an end-to-end pipeline for **dynamic ad object insertion**, combining video matching, surface detection, and compositing with latency comparable to video transcoding, serving as a proof of concept for online deployment.
- Accelerated SAM3 surface detection by **8.5**× via **adaptive frame sampling** and fine-tuned a Qwen3-VL embedding model for **brand-conflict detection and contextual matching**.

SELECTED PUBLICATIONS

[1] **Lingzhi Zhao**, Yu Liu, Yilin Li, Diane Xu, Shi Han Chong, Yang Fan Mia, Xinyi Liu, Yufei Li, Xinheng Sheng, Isaac Zhang, Chen Zheng, Zhengyuan Yang, Lei Shi, and Feng Qian. Short UGC Video Quality Assessment via Large Language Model: Benchmark and Method. ***Submitted to ICLR 2027.***

[2] **Lingzhi Zhao**, Yongqiang Gui, Yanyan Suo, Sandesh Dhawaskar Sathyanarayana, Mingyuan Wu, Bo Chen, Ruixiao Zhang, Klara Nahrstedt, and Shu Shi. MARS: Streaming How You Move with Motion-Aware ABR for Real-Time VR. ***Submitted to ACM MMSys 2027.***

[3] Beitong Tian*, **Lingzhi Zhao***, Bo Chen, Haozhen Zheng, Jingcheng Yang, Mingyuan Wu, Deepak Vasisht, and Klara Nahrstedt. AquaVLM: Improving Underwater Situation Awareness with Mobile Vision Language Models. (*: equal contributions) ***ACM UbiComp 2027.*** *(corresponding author)*

[4] Beitong Tian*, **Lingzhi Zhao***, Bo Chen, Mingyuan Wu, Haozhen Zheng, Deepak Vasisht, Francis Y. Yan, and Klara Nahrstedt. AquaScope: Reliable Underwater Image Transmission on Mobile Devices. (*: equal contributions) ***ACM MobiSys 2026.*** *(corresponding author)*

[5] **Lingzhi Zhao**, Yongqiang Gui, Yanyan Suo, Sandesh Dhawaskar Sathyanarayana, Ruixiao Zhang, Shu Shi, and Klara Nahrstedt. Trinity: Exploiting Latency Sensitivity to Improve Quality of Experience on Cloud VR Gaming. ***ACM MMSys 2026.***

[6] **Lingzhi Zhao**, Ying Cui, Yuhang Jia, Yunfei Zhang, and Klara Nahrstedt. Enhancing Neural Adaptive Wireless Video Streaming via Lower-Layer Information Exposure and Online Tuning. ***IEEE Transactions on Multimedia 2025.***

[7] **Lingzhi Zhao**, Qian Zhou, Bo Chen, and Klara Nahrstedt. 360LiveCast: A Low-Latency and Bandwidth-efficient Multicast Framework for Live 360 Video. ***IEEE MIPR 2025.***

[8] **Lingzhi Zhao**, Ying Cui, Sheng Yang, and Shlomo Shamai (Shitz). An Optimization Framework for General Rate Splitting for General Multicast. ***IEEE Transactions on Wireless Communications 2022.***

[9] **Lingzhi Zhao**, Ying Cui, Zhi Liu, Yunfei Zhang, and Sheng Yang. Adaptive Streaming of 360 Videos with Perfect, Imperfect, and Unknown FoV Viewing Probabilities in Wireless Networks. ***IEEE Transactions on Image Processing 2021.***